

Lecture 7: Markov Chain Monte Carlo

Professor Alexander Franks

2025-11-05

Reminders

- Quiz 3 today / normal, mcmc
- Hw 3 due, HW4 out tonight

$$y_1, \dots, y_n \stackrel{\text{iid}}{\sim} N(\mu, \sigma^2)$$

$$\mu \sim N(\mu_0, \frac{\sigma^2}{K_0})$$

$$P(\mu | y_1, \dots, y_n) \sim N(w\bar{y} + (1-w)\mu_0, \frac{\sigma^2}{n+K_0})$$

$$w = \frac{n}{n+K_0}$$

Bias - Var. Tradeoff
 $w\bar{y} + (1-w)\mu_0$
(random)

σ^2 also unknown

$$P(\mu, \sigma^2 | y_1, \dots, y_n) \propto \frac{1}{\sigma^2}$$

$$\underbrace{\mathcal{L}(\mu, \sigma^2)} \underbrace{P(\mu | \sigma^2)} \underbrace{P(\sigma^2)}$$

$$\nwarrow N(\mu_0, \frac{\sigma^2}{K_0})$$

Precision
 $\frac{1}{\sigma}$

Monte Carlo



Monte Carlo estimation

LLN

- $\bar{\theta} = \sum_{s=1}^S \theta^{(s)} / S \rightarrow \mathbb{E}[\theta | y_1, \dots, y_n]$

- $\sum_{s=1}^S (\theta^{(s)} - \bar{\theta})^2 / (S - 1) \rightarrow \text{Var}[\theta | y_1, \dots, y_n]$

- $\# (\theta^{(s)} \leq c) / S \rightarrow \Pr(\theta \leq c | y_1, \dots, y_n)$

- the α -percentile of $\{\theta^{(1)}, \dots, \theta^{(S)}\} \rightarrow \theta_\alpha$

Can write down posterior, but
we want summaries (integrals)

$\mathcal{L}(\theta) P(\theta)$

Sampling from the posterior distributions

- The Monte Carlo methods we discussed previously assumed we could easily get samples from the posterior, e.g. with `rnorm`
- In general, sampling from a general probability distribution is hard
- Want to call `rcomplicateddistribution()` but don't have it
 - Inversion sampling is limited, *Rejection sampling.*
 - ~~Grid sampling is reasonable in 1 or 2 dimensions~~
- In high dimensions, these approaches aren't sufficient

M C M C

Markov Chain Monte Carlo

- We want independent random samples, $\theta^{(s)}$ from $p(\theta \mid y_1, \dots, y_n)$
- But there is no good way to get independent samples
- Alternative, create a sequence of correlated samples that converge to the correct distribution
- Markov Chain Monte Carlo gives us a way to generate correlated samples from a distribution

Too Hard

We'll
pay a
price.

converge to the posterior.

This is the price we pay.

Monte Carlo Error

$$\text{Var}\left(\sum_{i=1}^S y_i\right) = \sum_{i,j} \text{Cov}(y_i, y_j) \quad \text{iff indep.} \\ = \sum \text{Var}(y_i)$$

- Reminder: $\bar{\theta} = \sum_{s=1}^S \theta^{(s)} / S$ and S is the number of samples.

- If the samples are independent,

M.C. Error

$$\text{Var}(\bar{\theta}) = \frac{1}{S^2} \sum_{s=1}^S \text{Var}(\theta^{(s)}) = \frac{\text{Var}(\theta \mid y_1, \dots, y_n)}{S}$$

- If the samples are positively correlated, (as with MCMC)

$$\text{Var}(\bar{\theta}) = \frac{1}{S^2} \sum_{s,t} \text{Cov}(\theta^{(s)}, \theta^{(t)}) > \frac{\text{Var}(\theta \mid y_1, \dots, y_n)}{S}$$

- MCMC methods have higher Monte Carlo error due to positive dependence between samples.
- Hope to minimize dependence and thus MC error

MCMC error higher than MC 5 / 86

because of the ' correlation.

Basics of Markov Chains

Markov Chains: Big Picture

LLN

- For standard Monte Carlo, we make use of the law of large number to approximate posterior quantities
- The law of large numbers can still apply to random variables that are not independent
- We have a sequence of random variables indexed in "time", θ_t
- We'll be using a discrete-time Markov Chain: $t \in 0, 1, \dots, T$
- The observations, $\theta^{(t)}$ can be discrete or continuous ("discrete-state" or "continuous-state" Markov Chain)

Discrete-state Markov Chains

- Let $\theta^{(t)} \in 1, 2, \dots, M$ be the state space for the Markov Chain
- A sequence is called a Markov chain if


$$Pr(\theta^{(t+1)} \mid \theta^{(t)}, \cancel{\theta^{(t-1)} \dots \theta^{(1)}}) = Pr(\theta^{(t+1)} \mid \theta^{(t)})$$

for all $t \geq 0$

- **The Markov property:** given the entire past history, $\theta^{(1)}, \dots, \theta^{(t)}$, the most recent $\theta^{(t+1)}$ depends only on the immediate past, $\theta^{(t)}$

The Transition Matrix

- Define $q_{ij} = Pr(\theta^{(t+1)} | \theta^{(t)})$ is the transition probability from state i to state j
TO From
- The $M \times M$ matrix $Q = (q_{ij})$ is called the *transition matrix* of the Markov Chain



3-state example

$$Q = \begin{matrix} & \text{TO} & \\ \text{From} & \begin{bmatrix} q_{11} & q_{12} & q_{13} \\ q_{21} & q_{22} & q_{23} \\ q_{31} & q_{32} & q_{33} \end{bmatrix} \end{matrix}$$

The Transition Matrix

3-state example

$$Q = \begin{bmatrix} q_{11} & q_{12} & q_{13} \\ q_{21} & q_{22} & q_{23} \\ q_{31} & q_{32} & q_{33} \end{bmatrix}$$

- The rows of the transition matrix sum to 1
- Note: $Q^n = (q_{ij}^{(n)})$ is the probability of transitioning from i to j in n steps

A handwritten red diagram illustrating the calculation of Q^n . It shows the expression $Q \times Q \times \dots \times Q$ with Q at both ends and an ellipsis in the middle. A red bracket underneath the entire sequence of Q 's is labeled with a red n . A red arrow points from the word "steps" in the list item above to the first Q in the sequence.

The limiting distribution

- A regular, irreducible Markov chain has a **limiting probability distribution**
 - Cover definitions of regular and irreducible in PSTAT160 (or related)
- Limit distribution describes the long-run fraction of time the Markov Chain spends in each state
 - *Does not* depend on where the chain starts
- Let $\pi = (\pi_1, \dots, \pi_M)$ be a row vector of probabilities associated with each state, such that $\sum_{i=1}^M \pi_i = 1$
 - The limiting distribution converges to π , which is said to be **stationary** because $\pi Q = \pi$
 - If you sample from the limiting distribution and then transition, the result is still distributed according to the limiting distribution

Markov Chain Example

- Sociologists often study social mobility using a Markov chain.
- In this example, the state space is {low income, middle income, and high income} of families
- Let $\mathbf{Q}$ be the transition matrix from parents income to childrens income

		<i>Kids</i>		
		Lower	Middle	Upper
$\mathbf{Q} =$ <i>Parents</i>	Lower	0.40	0.50	0.10
	Middle	0.05	0.70	0.25
	Upper	0.05	0.50	0.45

Multi-step Transition Probabilities

2-step transition probabilities

Grand kids

Parents $Q^2 = Q \times Q =$

0.1900	0.6000	0.2100
0.0675	<u>0.6400</u>	<u>0.2925</u>
0.0675	0.6000	0.3325

4-step transition probabilities

$$Q^4 = Q^2 \times Q^2 = \begin{vmatrix} 0.0908 & 0.6240 & 0.2852 \\ 0.0758 & 0.6256 & 0.2986 \\ 0.0758 & 0.6240 & 0.3002 \end{vmatrix}$$

Multi-step Transition Probabilities

4-step transition probabilities

$$\mathbf{Q}^4 = \mathbf{Q}^2 \times \mathbf{Q}^2 = \begin{vmatrix} 0.0908 & 0.6240 & 0.2852 \\ 0.0758 & 0.6256 & 0.2986 \\ 0.0758 & 0.6240 & 0.3002 \end{vmatrix}$$

8-step transition probabilities

$$\mathbf{Q}^8 = \mathbf{Q}^4 \times \mathbf{Q}^4 = \begin{vmatrix} 0.0772 & 0.6250 & 0.2978 \\ 0.0769 & 0.6250 & 0.2981 \\ 0.0769 & 0.6250 & 0.2981 \end{vmatrix}$$

Stationary distribution:

$$\begin{pmatrix} \pi_1 & \pi_2 & \pi_3 \\ \pi_1 & \pi_2 & \pi_3 \\ \pi_1 & \pi_2 & \pi_3 \end{pmatrix} = \lim_{n \rightarrow \infty} \mathbf{Q}^n$$

$\pi = (\pi_1, \pi_2, \pi_3)$
 $\pi_1 + \pi_2 + \pi_3 = 1$

The limiting distribution

$$\mathbf{Q}^\infty = \mathbf{1}\pi = \begin{vmatrix} \pi_1 & \pi_2 & \pi_3 \\ \pi_1 & \pi_2 & \pi_3 \\ \pi_1 & \pi_2 & \pi_3 \end{vmatrix}$$

```
Q <- matrix(c(0.4, 0.05, 0.05,  
              0.5, 0.7, 0.5,  
              0.1, 0.25, 0.45),  
            ncol=3)
```

```
p <- eigen(t(Q))$vectors[, 1]  
stationary_probs <- p/sum(p)  
stationary_probs
```

↑ is related
to eigenvectors
of Q.

π ## [1] 0.07692308 0.62500000 0.29807692

```
stationary_probs %*% Q
```

```
##      [,1] [,2] [,3]  
## [1,] 0.07692308 0.625 0.2980769
```

Get π back.

Markov Chain Monte Carlo

- Incredible idea: create a Markov Chain with the desired limiting distribution

- Want the limiting distribution to be the posterior distribution

- Unlike the previous examples, we will mostly work with *infinite* state space

- Want $p(\theta^{(t+1)} | \theta^{(t)})$ to have limiting distribution $p(\theta | y)$

- If we run the random walk for long enough, $\theta^{(t)}$ will be distributed approximately according to $p(\theta | y)$

← Transition Rule

Analog of Q , transition matrix,
for continuous states.

"Transition kernel"

Generalizing the Rejection Sampler.

- Sample from proposal $q(\theta)$
 - Doesn't need to envelope $P(\theta|y)$!
 - If $P(\theta|y) > 0$, then $q(\theta) > 0$
- Unlike the rejection sampler, we never "throw out" samples.
- At each iteration, choice:
 - Accept the proposed sample
 - Or Repeat the last sample.

The Independence Sampler

- The Metropolis algorithm tells us how to construct a transition matrix with the correct limiting distribution
 - The Independence Sampler is a special case of the Metropolis algorithm
- Sample from a proposal, $q(\theta)$. Best if $q(\theta)$ is close to $p(\theta | y)$.
- If $p(\theta | y) > 0$ then we need $q(\theta) > 0$
- At each iteration we have a choice:
 - Accept the new proposed sample
 - Or keep the previous sample for another iteration

The Independence Sampler (modification of R-S.

1. Initialize θ_0 to be the starting point for you Markov Chain

2. Choose a *proposal distribution*, $J(\theta^*)$ (same as q before)

- Propose a candidate value for the next sample
- Best performance if density is very similar to target

3. Generate the candidate θ^* from the proposal distribution, J


4. Compute $r = \min(1, \frac{p(\theta^*|y)}{p(\theta_t|y)})$ "ratio of posterior evaluated at proposed div. by current"

5. Set $\theta_{t+1} \leftarrow \theta^*$ with probability r

- Generate a uniform random number $u \sim \text{Unif}(0, 1)$
- If $u < r$ we accept θ^* as our next sample
- ~~Else $\theta_{t+1} \leftarrow \theta_t$ (we do not update the sample this time)~~

Else set $\theta_{t+1} \leftarrow \theta_t$

Intuition




- If $p(\theta^* | y) > p(\theta_t | y)$ accept with probability 1
 - The proposed sample has higher posterior density than the previous sample
 - Always accept if we increase the posterior probability density
- If $p(\theta^* | y) < p(\theta_t | y)$ accept with probability $r < 1$
 - Accept with probability less than 1 if probability density would decrease
 - Relative frequency of θ^* vs θ_t in our samples should be $\frac{p(\theta^* | y)}{p(\theta_t | y)}$

An Example


- Let $P(\theta \mid y)$ be a Beta(5, 10) posterior distribution
- Propose from a distribution $J(\theta^*) \sim N(0.5, 1)$

```
## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.  
## i Please use `linewidth` instead.  
## This warning is displayed once every 8 hours.  
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was  
## generated.
```

```
## Warning: A numeric `legend.position` argument in `theme()` was deprecated  
## 3.5.0.  
## i Please use the `legend.position.inside` argument of `theme()` instead.  
## This warning is displayed once every 8 hours.  
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was  
## generated.
```




Independence Sampler



Note and source of confusion: samples are correlated over time for the "independence sampler".


Weighting by waiting



Where did the sampler get stuck? Where does it quickly leave?

Independence Sampler

Monte Carlo vs True




N = 100000 Bandwidth = 0.01065

The Metropolis Algorithm

- Generalize the previous special case
- Allow the proposal distribution to depend on the most recent sample
 - Sometimes called an "Independence sampler":
 $J(\theta^*)$, e.g. $\theta^* \sim N(0.5, 1)$
 - **Metropolis**: $J(\theta^* | \theta_t)$, e.g. $\theta^* \sim N(\theta_t, 1)$
- Independence sampler: "Independence" refers to the proposal being fixed (the samples are **not** independent)!
- **Metropolis sampler: a "moving" proposal distribution**

Propose new
parameters
given
current.

The Metropolis Algorithm



1. Initialize θ_0 to be the starting point for your Markov Chain

2. Choose a proposal distribution, $J(\theta^* | \theta_t)$, e.g. $\theta^* \sim \mathcal{N}(\theta_t, 1)$

- Propose a candidate value for the next sample

- Must have symmetry: $J(\theta^* | \theta_t) = J(\theta_t | \theta^*)$

3. Generate the candidate θ^* from the proposal distribution, J

4. Compute $r = \min(1, \frac{p(\theta^*|y)}{p(\theta_t|y)})$ same rule

5. Set $\theta_{t+1} \leftarrow \theta^*$ with probability r

- Generate a uniform random number $u \sim \text{Unif}(0, 1)$

- If $u < r$ we accept θ^* as our next sample

- Else $\theta_{t+1} \leftarrow \theta_t$ (we do not update the sample this time)

Else keep previous, $\theta_{t+1} \leftarrow \theta_t$

Equally likely
to go from
 $\theta_t \rightarrow \theta^*$ as
 $\theta^* \rightarrow \theta_t$

Aside: Arianna Rosenbluth



https://en.wikipedia.org/wiki/Arianna_W._Rosenbluth

Metropolis Algorithm

- Let $P(\theta \mid y)$ be a Beta(5, 10) posterior distribution
- 1-d sampling: lets try sampling from the Beta using the Metropolis algorithm
- Initialize θ_0 to 0.9
 - Note that the probability of drawing a value larger than 0.9 from a Beta(5, 10) is smaller than $1e-8$
 - Our initial value is far from the high posterior density
 - In the long run this won't matter

- Define transition kernel $J(\theta_{t+1} \mid \theta_t)$ as $\theta^* \sim N(\theta_t, \tau^2)$
 - How does choice of τ^2 effect performance of MC sampler?

How do
I choose
 τ ?


ϵ^2 : not too big,
not too small,
Goldilocks problem.

Demo

Autocorrelation of the Markov Chain

$\tau^2 = 0.01$ ("small" proposal variance)

Plot θ^t vs θ^{t+5} for all values of t




Correlation = 0.9643696

Autocorrelation of the Markov Chain

$\tau^2 = 0.01$ ("small" jump)

Plot θ^t vs θ^{t+100} for all values of t




The Metropolis Algorithm

$\tau^2 = 0.01$ ("small" jump)

```
acf(rw_small, lag=100)
```

Auto-correlation Function
(acf)



Want correlations to be small (independence)

```
print(sprintf("Effective sample size: %.2f, Rejection Rate: %.2f",  
             effectiveSize(rw_small), rejectionRate(as.mcmc(rw_small))))
```


```
## [1] "Effective sample size: 4.42, Rejection Rate: 0.05"
```

The Metropolis Algorithm

$\tau^2 = 0.01$ ("small" jump)

*Run MCMC for 1000
iters.*

samples vs target




*Terrible
Error.*

N = 1000 Bandwidth = 0.01917

The Metropolis Algorithm




$\tau^2 = 0.1$ ("medium" proposal variance)



```
acf(rw_mid, lag=50)
```




```
print(sprintf("Effective sample size: %.2f, Rejection Rate: %.2f",  
             effectiveSize(rw_mid), rejectionRate(as.mcmc(rw_mid
```

```
## [1] "Effective sample size: 128.86, Rejection Rate: 0.23"
```



Less correlation.




time

The Metropolis Algorithm

$\tau^2 = 0.1$ ("medium" proposal variance)

samples vs target




N = 1000 Bandwidth = 0.02408
(moderate proposal variance)

The Metropolis Algorithm

$\tau^2 = 2$ ("large" jump)

```
acf(rw_large, lag=50)
```




```
print(sprintf("Effective sample size: %.2f, Rejection Rate: %.2f",  
             effectiveSize(rw_large), rejectionRate(as.mcmc(rw_large))))
```

```
## [1] "Effective sample size: 32.51, Rejection Rate: 0.93"
```

The Metropolis Algorithm


$\tau^2 = 2$ ("large" proposal variance)



N = 1000 Bandwidth = 0.02113
(large proposal variance)

Trace Plot:

Plot of parameter(s) vs
time (iteration)





Small, Moderate and Large Proposal variance

"City Scape" $\Rightarrow$ High Rejection Rate
 traceplot

$J(\theta^* | \theta_t)$
 $\sim N(\theta_t, \sigma^2)$

σ^2 big





True Posterior

MCMC approx.

Lots of correlation



No "flat spots", High Acceptance.

σ^2 small




Lots of correlation.

σ^2 "goldilocks"



10,000 Samples



[1] "Effective sample size (small proposal variance): 20.68"

[1] "Effective sample size (medium proposal variance): 1041.74"

[1] "Effective sample size (large proposal variance): 418.34"

MCMC diagnostics

- Diagnose the chain performance by examining:

- Rejection rate
- Autocorrelation
- Effective sample size

• Trace Plot (visual inspection)

MCMC diagnostics

- Too High*
- **Rejection rate:** if rejection rate is high, the proposal density is proposing "too far away" from the current sample
 - Means we are often keeping the previous sample
 - Sampler is "sticky". Traceplots look like cityscapes.

Fraction of proposals that aren't used.

Too Low

- Accept too much, not exploring parameter space fast enough.


MCMC diagnostics

- **Autocorrelation:** if samples are highly correlated, the proposal density is proposing "too close" to the current sample
 - Highly correlated implies the Markov chain is mixing slowly
- The mixing time of a Markov chain is the time until the Markov chain is "close" to its limiting distribution.

$\Rightarrow$ High
mc
error.

$$\text{Cor}(\theta_t, \theta_{t+l})$$

'acf' function



MCMC diagnostics

- **Effective sample size:** correlated chain of samples is equivalent to this number of independent samples *(in terms of MC error)*
 - High rejection rate implies a lot of duplicate samples so effective size is smaller than number of iterations
 - High autocorrelation means neighboring samples are very similar (even if not exactly the same)

For independent samples, MC error is
proportional to $1/\sqrt{ESS}$

For dep. samples:

$$= \frac{\sum \sum \text{cov}(\theta_i, \theta_j | y)}{n^2}$$

Solve For θ_{ESS}

Metropolis Algorithm

- In the Beta example, the accept ratio, $\min(1, \frac{p(\theta^*|y)}{p(\theta_t|y)})$ is zero when $\theta^* > 1$ or $\theta^* < 0$
- If τ^2 (proposal variance) too large, you will often reject your proposal
 - Proposing far from your current location may move you too far out of the high density areas
 - This makes for a "sticky" chain (stay at current sample for a long time)
- τ^2 too small, the chain explore the parameter space slowly
- A rule of thumb is to aim for 30%-40% acceptance rate for random walk samplers.
 - This balances "stickiness" and slow convergence

"City Sega"

high acc,

Metropolis-Hastings Algorithm

- The Metropolis-Hastings algorithm allows us to use non-symmetric proposals
- The Hastings correction is needed when $J(\theta^* | \theta_t) \neq J(\theta^t | \theta^*)$

$$r = \min\left(1, \frac{p(\theta^* | y)}{p(\theta_t | y)} \frac{J(\theta^t | \theta^*)}{J(\theta^* | \theta_t)}\right)$$

- For symmetric proposals $\frac{J(\theta^t | \theta^*)}{J(\theta^* | \theta_t)} = 1$

Hastings correction.

MCMC for multivariate distributions

- Modeling wing length of different species of midge (small, two-winged flies)
- Reminder: $Y_i \sim N(\mu, \sigma^2)$
- $P(\mu \mid \sigma^2), \mu \sim N(\mu_0, \frac{\sigma^2}{\kappa_0})$
- $P(\sigma^2) \propto \frac{1}{\sigma^2}$ (improper prior)


MCMC for multivariate distributions

Example: midge wing length


- Modeling wing length of different species of midge (small, two-winged flies)
- From prior studies: mean wing length close to 1.9mm with sd close to 0.1mm
- $\mu_0 = 1.9, \sigma_0^2 = 0.01$
- Choose $\kappa_0 = 1$
- We will run 2 separate chains at different starting locations
- $J(\mu^*, \log(\sigma^*)) \sim N((\mu^t, \log(\sigma^t)), \begin{pmatrix} \tau_\mu^2 & 0 \\ 0 & \tau_{\log\sigma}^2 \end{pmatrix})$

Initialization and Convergence

- In the long run, it doesn't matter where you initialize your sampler
- In practice, we can only run an algorithm for a finite amount of time
 - Need to check with the sampler has converged to the limiting distribution
 - Exclude samples close to the initial values since these are unlikely to be representative samples
 - We call the time to convergence burn in and throw away and samples generated during this time.
- How do we know when a sampler has converged to the limiting distribution?



Run multiple Chains in Parallel.
(Different Initializations)



Running multiple chains

- How do we know when a sampler has converged to the limiting distribution?
 - Hard to know for sure.
- Idea: run multiple chains at very different initial locations
 - If the chains are very far apart we haven't converged
 - If the chains end up in ¹the same place¹¹, we have confidence that its reached convergence

Diagnosing mixing with Rhat

$$B = \frac{n}{m-1} \sum_{j=1}^m \left(\bar{\psi}_{.j} - \bar{\psi}_{..} \right)^2$$

$$W = \frac{1}{m} \sum_{j=1}^m s_j^2$$

Demo

Multiple Chains in Stan

```
library(cmdstanr)
y <- c(1.64, 1.7, 1.72, 1.74, 1.82, 1.82, 1.82, 1.9, 2.08)

# compile stan model. This may take a minute.
stan_model <- cmdstan_model(stan_file="normal_model.stan")

## data is a list, arguments must match the datablock in the stan
stan_fit <- stan_model$sample(data=list(N=length(y), y=y, k0=1),
                              chains = 2,
                              num_warmup=200,
                              num_samples = 1000,
                              refresh=200)
```

Printing,

Burnin

keep.

2 x 1000 total Samples.

Multiple Chains in Stan

```
## Running MCMC with 2 sequential chains...
##
## Chain 1 finished in 0.0 seconds.
## Chain 2 finished in 0.0 seconds.
##
## Both chains finished successfully.
## Mean chain execution time: 0.0 seconds.
## Total execution time: 0.3 seconds.
```

```
SAMPLING FOR MODEL 'normal_model' NOW (CHAIN 1).
Chain 1:
Chain 1: Gradient evaluation took 9e-06 seconds
Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.09 seconds.
Chain 1: Adjust your expectations accordingly!
Chain 1:
Chain 1:
Chain 1: Iteration: 1 / 1000 [ 0%] (Warmup)
Chain 1: Iteration: 200 / 1000 [ 20%] (Warmup)
Chain 1: Iteration: 201 / 1000 [ 20%] (Sampling)
Chain 1: Iteration: 400 / 1000 [ 40%] (Sampling)
Chain 1: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 1:
Chain 1: Elapsed Time: 0.015009 seconds (Warm-up)
Chain 1: 0.02576 seconds (Sampling)
Chain 1: 0.040769 seconds (Total)
Chain 1:
```

```
SAMPLING FOR MODEL 'normal_model' NOW (CHAIN 2).
Chain 2:
Chain 2: Gradient evaluation took 1.5e-05 seconds
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.15 seconds.
Chain 2: Adjust your expectations accordingly!
Chain 2:
Chain 2:
Chain 2: Iteration: 1 / 1000 [ 0%] (Warmup)
Chain 2: Iteration: 200 / 1000 [ 20%] (Warmup)
Chain 2: Iteration: 201 / 1000 [ 20%] (Sampling)
Chain 2: Iteration: 400 / 1000 [ 40%] (Sampling)
Chain 2: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 2: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 2:
```

Stan Samples

```
draws <- stan_fit$draws(format="df")  
draws
```


log P(b|y)

```
## # A draws_df: 1000 iterations, 2 chains, and 3 variables  
##   lp__    mu  sigma  
## 1    17  1.9  0.16  
## 2    17  1.8  0.15  
## 3    18  1.8  0.14  
## 4    17  1.8  0.16  
## 5    18  1.8  0.11  
## 6    18  1.8  0.12  
## 7    18  1.8  0.12  
## 8    18  1.8  0.11  
## 9    16  1.8  0.20  
## 10   17  1.9  0.15  
## # ... with 1990 more draws  
## # ... hidden reserved variables {'.chain', '.iteration', '.draw'}
```


Plotting Stan Draws


```
mcmc_hist(draws, pars=c("mu", "sigma"))
```

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.




Plotting Stan Draws

```
mcmc_scatter(draws,  
             pars=c("mu", "sigma"),  
             transformations=list(sigma=function(x) 1/x^2))
```




MCMC Diagnostics: Traceplots

```
mcmc_trace(draws, pars=c("mu", "sigma"))
```



MCMC Diagnostics: Traceplots

```
mcmc_trace(draws, pars=c("mu"), window = c(1, 100))
```



MCMC Diagnostics

Efficiency can be expressed as (effective samples) / (total iterations)

```
stan_fit$summary() %>% glimpse
```


```
## Rows: 3
## Columns: 10
## $ variable <chr> "lp_", "mu", "sigma"
## $ mean <dbl> 17.2356446, 1.8138022, 0.1308553
## $ median <dbl> 17.5798500, 1.8129550, 0.1251045
## $ sd <dbl> 1.12399902, 0.04227812, 0.03454343
## $ mad <dbl> 0.82136040, 0.03759132, 0.02915681
## $ q5 <dbl> 15.01483500, 1.74625850, 0.08693568
## $ q95 <dbl> 18.3126400, 1.8858515, 0.1942028
## $ rhat <dbl> 1.000395, 1.000797, 1.000443
## $ ess_bulk <dbl> 843.4363, 1394.9833, 778.1124
## $ ess_tail <dbl> 937.1279, 1085.1102, 757.5916
```

Post.
Mean

Post.
Quantiles

Effective S.S.

MCMC for multivariate distributions



MCMC for multivariate distributions

- We will run 2 separate chains at different starting locations


- For chain 1: $J(\theta^* \mid \theta_t) = N \left(\begin{pmatrix} \mu^{(t)} \\ \log(\sigma^{(t)}) \end{pmatrix}, \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} \right)$

- For chain 2: $J(\theta^* \mid \theta_t) = N \left(\begin{pmatrix} \mu^{(t)} \\ \log(\sigma^{(t)}) \end{pmatrix}, \begin{pmatrix} 0.1 & 0 \\ 0 & 0.1 \end{pmatrix} \right)$


How to
choose?

(μ_t, σ_t)

MCMC for multivariate distributions



MCMC for multivariate distributions



```
## # A tibble: 2 × 3
##   chain    EF_mu EF_prec
##   <fct>    <dbl>   <dbl>
## 1 chain1  389.    108.
## 2 chain2   7.43    35.6
```

Metropolis Sampling

Try out the Metropolis algorithm at:

<https://chi-feng.github.io/mcmc-demo/app.html>

- Choose "Random Walk MH" algorithm
- Experiment by sampling from different target distribution
- Try different proposal variances by changing "Proposal σ "

Gibbs Sampling

The Gibbs Sampler

- The Gibbs sampler is actually special case of the MH sampler
 - This is not obvious or immediately apparent
- Idea: break the problem down into many smaller sampling problems
- Iteratively update each parameter in the full parameter vector by doing getting lower-dimensional sample

The Gibbs Sampler

- Suppose that the parameter vector θ can be divided into d subvectors.
- For iterations, $s = 1, \dots, S$:
 - For $j \in 1, \dots, d$
 - Draw a value from the conditional distribution of θ_j given all the other parameters,

$$\theta_j^s \sim p(\theta_j | \theta_{-j}^{s-1}, y),$$

where θ_{-j}^{t-1} has consists of updated parameters for all parameters preceding j and the previous iteration's values for all succeeding parameters,

$$\theta_{-j}^{t-1} = \left(\theta_1^t, \dots, \theta_{j-1}^t, \theta_{j-1}^{t-1}, \dots, \theta_d^{t-1} \right).$$

The Gibbs Sampler

- To identify the full conditional distributions:
 1. Write down the full posterior
 2. For each parameter, θ_j , remove all multiplicative constants that don't have a θ_j in them.
 3. Identify the type of distribution to sample from
- Assuming conjugate prior distributions, the full conditionals in the normal model are:
 - $p(\mu \mid \sigma^2, y)$ is a normal distribution
 - $p(\frac{1}{\sigma^2} \mid \mu, y)$ is a gama distribution

The Gibbs Sampler



Gibbs Sampling a Bivariate Normal

Demo

Gibbs Sampler

- Advantages
 - proposals, $p(\theta^{(t+1)} \mid \theta^{(t)})$, are never rejected.
 - don't need to choose the proposal density distribution or tune parameters of the density
- Disadvantages
 - It can be difficult to derive the full-conditional distributions (unless, for example, the prior distributions are chosen to be conjugate)
 - When the parameters of the posterior distribution are highly correlated:
 - Hard to traverse diagonals and consequently:
 - autocorrelation of the samples is high
 - effective sample size is low

Challenges in MCMC

- Modern models often have *many* parameters. Large models pose a challenge for MCMC.
- When there are thousands or more parameters
 - MCMC may take a long time to converge to the limiting distribution
 - In Metropolis-Hastings we have many tuning parameters for the proposal distribution
 - Gibbs sampling has no tuning parameters, but does not work well for highly correlated posterior distributions
- In general, MCMC is slow relative to optimization methods

Summary of MCMC

What is Monte Carlo and MCMC:

- MCMC is *not* a model
- It does *not* generate more information
- Provides *dependent* approximate samples from $p(\theta \mid y)$
- Samples can be used to summarise $p(\theta \mid y)$ (approximate integrals)

Computational Considerations

- For very small values of $p(\theta \mid y)$, numerical underflow is a problem
- Can resolve this by working on the log scale

```
dnorm(1000) / dnorm(1001)
```

```
## [1] NaN
```

```
dnorm(1000, log=TRUE) - dnorm(1001, log=TRUE)
```

```
## [1] 1000.5
```

- Compute $l = \min(0, \log(p(\theta^* \mid y)) - \log(p(\theta_t \mid y)))$
- If $\log(u) < l$ we accept θ^* as our next point

MCMC for multivariate distributions

- Modeling wing length of different species of midge (small, two-winged flies)
- Reminder: $Y_i \sim N(\mu, \sigma^2)$
- $P(\mu \mid \sigma^2), \mu \sim N(\mu_0, \frac{\sigma^2}{\kappa_0})$
- $P(\sigma^2), \sigma^2 \sim \text{Inv-Gamma}(\nu_0/2, \nu_0/2\sigma_0^2)$
- Parameterize in terms of $\theta = (\mu, \log(\sigma))$? Why parameterize this way?
- Let $J(\theta^* \mid \theta_t) = N \left(\begin{pmatrix} \mu^{(t)} \\ \log(\sigma^{(t)}) \end{pmatrix}, \begin{pmatrix} 0.5 & 0 \\ 0 & 0.5 \end{pmatrix} \right)$

Challenges in MCMC

- Modern models often have *many* parameters. Large models pose a challenge for MCMC.
- When there are thousands or more parameters
 - MCMC may take a long time to converge to the stationary distribution
 - In Metropolis-Hastings we have many tuning parameters for the proposal distribution
 - Gibbs sampling has no tuning parameters, but does not work well for highly correlated posterior distributions
- In general, MCMC is very slow relative to optimization methods


Modern MCMC

- Gibbs and Metropolis samplers have a "random walk" behavior
 - Induces autocorrelation
 - Makes it difficult to explore the posterior space
- Hamiltonian Monte Carlo (HMC) borrows an idea from physics to reduce this problem

HMC (Hamiltonian Monte Carlo)

- Imagine a marble on a frictionless surface. The location of the marble is the current value of θ_t
- The negative posterior density is the "height" of the surface
- Each iteration we flick the marble with some velocity in a random direction
- Regions of high posterior density are like "wells"

Stan
uses
it.



HMC

- For Metropolis-Hastings we only need to be able to evaluate the posterior at each location

$$P(\theta | z)$$

- For HMC we need the gradient (derivative) of the posterior as well

- Determines where the marble rolls

$$\frac{dP(\theta | y)}{d\theta}$$

- In physics the Hamiltonian is the sum of the kinetic energies, plus the potential energy of the particles

- As our proposal, we randomly sample a momentum for the marble and update its position accordingly

- Can think of HMC as the MH algorithm with a very clever jumping/proposal rule

$$J(\theta^* | \theta_t) \text{ s.t.}$$

θ^* far from θ_t and we accept.

HMC

Try out HMC at:

<https://chi-feng.github.io/mcmc-demo/app.html>

- Choose "HamiltonianMC" algorithm
- Experiment by sampling from different target distributions
- Compare to the Random Walk Metropolis

Approximate Inference

- MCMC can be very slow in high dimensional problems
- Idea: find a distribution that is easy to sample from which closely approximate $p(\theta \mid y)$
- A couple of examples
 - Laplace Approximation
 - Variational Bayes

Laplace Approximation to the Posterior

- Approximate the posterior distribution using a multivariate normal distribution
- When we have a lot of i.i.d. observations, the posterior will be approximately normal
- Center the normal at the mode of the posterior
- Compute the (co)variance of the normal by computing the second derivative / hessian of the posterior at the mode

Laplace Approximation

- Approximate the posterior distribution using a normal distribution
- When we have a lot of i.i.d. observations, the posterior will be approximately normal
- Center the normal at the mode of the posterior
- Compute the (co)variance of the normal by computing the second derivative / hessian of the posterior at the mode

Laplace Approximation

- Let $\tilde{\theta}$ be the mode of the posterior distribution
- Use a Taylor Series approximation the log-posterior around the mode is
 - $\log P(\theta | y) \approx \log P(\tilde{\theta} | y) - 1/2(\theta - \tilde{\theta})H(\theta - \tilde{\theta})$
 - $H = \frac{d^2}{d\theta^2} \log p(\theta | y)$
 - Note, linear term falls out because derivative at the mode is zero
- $p(\theta | y) \approx N(\tilde{\theta}, I(\theta)^{-1})$

Finding the mode of the posterior distribution

- Calculus
 - Take the log
 - Differentiate, set to zero and solve
- Computational
 - `optim` in R for one dimensional posteriors
 - `optimise` in R for multivariate p

Variational Bayes

- Find even better approximations
- Let θ be d dimensional parameter vector
- Let $\epsilon \sim MVN_d(0, \Sigma)$
- Let g_λ be a class of flexible functions parameterized by λ
 - Neural networks!
- Solve an optimization problem:

$$\arg \min_{\lambda} \text{dist}(p(\theta \mid y), p(g_\lambda(\epsilon)))$$

- Minimize the "distance" between the true posterior and the approximate one.
- Sample ϵ from a multivariate normal. $g_\lambda(\epsilon)$ will be a sample from something close to $p(\theta \mid y)$